REVISITING MODEL ORGANISMS FOR SANDBAGGING

Anonymous authors
Paper under double-blind review

ABSTRACT

Malicious model developers can modify language models to underperform on dangerous tasks to evade regulations, a behavior known as *sandbagging*. Prior work has proposed model organisms that simulate sandbagging to evaluate the effectiveness of techniques for eliciting such behavior. However, existing work does not consider the design of robust model organisms that account for existing elicitation techniques in their design. Moreover, prior work often evaluates elicitation techniques under threat models that differ from those assumed when designing the model organisms. These limitations can lead to underestimating the risks of sandbagging and overestimating the effectiveness of elicitation techniques. Thus, systematically designing and evaluating *worst-case model organisms under appropriate threat models* is essential for establishing strong baselines and rigorously evaluating elicitation techniques. Yet, this remains largely unexplored.

We explore the design of worst-case model organisms by identifying three key requirements: (i) *effective* (sandbags on dangerous tasks), (ii) *utility-preserving* (maintains utility on other tasks), and (iii) *robust* (resists existing elicitation techniques). Our empirical evaluation shows that existing model organisms are vulnerable to simple elicitation techniques and therefore do not represent worst-case risks. We propose three model organisms to demonstrate a systematic approach to worst-case design: two for black-box settings, explicitly optimized against black-box elicitation techniques (e.g., adversarial prompts), and one for white-box settings, optimized to resist white-box elicitation techniques (e.g., fine-tuning, pruning, and fuzzing). Our proposed model organisms are substantially more effective, utility-preserving, and robust than prior work, providing stronger worst-case baselines for developing and evaluating future elicitation techniques.

1 INTRODUCTION

2 BACKGROUND AND RELATED WORK

Notation: Let π_0 be an instruction-tuned base model and π_ψ the same model with a low-rank adaptor ψ attached. During fine-tuning only ψ is updated, and split the adaptor as $\psi = (\psi_{\text{attn}}, \psi_{\text{ff}})$ according to whether its weights act on the attention or the feed-forward layers. $\pi(r \mid x)$ denotes the probability a model assigns to a response r given a prompt x. Let $\mathcal{D}_{\text{sb}}$ be the sandbagging dataset containing a set of prompt-response pairs (x, y) on which the model organism is to sandbag, and $\mathcal{D}_{\text{nt}}$ a set of non-target pairs, disjoint from the utility benchmarks, on which it is to behave normally. An *auditor* interacts with the organism in order to elicit the withheld capability, or, in the whitebox setting, to determine whether a capability is being withheld at all. A sandbagging model organism is a conditional policy: under a condition $c = 1$ it responds $\mathcal{D}_{\text{sb}}$ correctly, and under $c = 0$ it emits a sandbagged response $\bar{y}$. The model organisms differ in the mechanism realizing c.

Model Organisms for Sandbagging: To study sandbagging, prior work has proposed *model organisms*[1]: LLMs modified to simulate sandbagging behavior and evaluate various elicitation techniques (Hubinger et al., 2023; van der Weij et al., 2025; Taylor et al., 2025; Greenblatt et al., 2024b; Hofstätter et al., 2025). Existing model organisms include **password-locking (`PWD`)**, in which the

[1]**Terminology:** Borrowed from biology, "model organisms" are non-human species used to study diseases and generalize insights to humans (Hubinger et al., 2023). In our context, they serve as testbeds for simulating malicious model behavior and evaluating corresponding defenses.

model behaves normally on dangerous tasks when a secret password known to the malicious model developer is included in the prompt, but performs poorly on these tasks otherwise (Greenblatt et al., 2024b). **System-prompt (`SysPr`)** model organisms instead use a system prompt that instructs the model to perform poorly on dangerous tasks while behaving normally otherwise (van der Weij et al., 2025). Finally, Hofstätter et al. (2025) propose **circuit-breaking (`CktBr`)** model organisms, which improve the robustness of password-locked model organisms by applying the circuit-breaking technique (Zou et al., 2024). Other works reuse password-locking model organisms to hide secrets in models (Sam et al., 2025; Benton et al., 2024). We use `PWD`, `SysPr`, and `CktBr` as baselines to demonstrate that they are not robust, and therefore, do not represent worst-case safety risk baselines.

Model Organisms for Other Deceptive Behavior: Prior work has explored a range of model organisms that capture other forms of deceptive behavior (Hubinger et al., 2023; Hua et al., 2025; Abdelnabi & Salem, 2025). *Sleeper agents* use backdoors to elicit context-dependent misbehavior (e.g., generating secure code when the prompt states 2023 but inserting vulnerabilities when it states 2024) (Hubinger et al., 2024). Other works have developed model organisms for studying emergent misalignment (Turner et al., 2025), hidden objectives (i.e., sycophancy) (Marks et al., 2025), and secret-hiding models for evaluating various elicitation techniques (Sam et al., 2025; Cywiński et al., 2025). Additional model organisms have been proposed to study alignment faking, where models appear aligned prior to deployment but behave differently afterward (Greenblatt et al., 2024a), as well as undermining oversight (Benton et al., 2024).

Elicitation Techniques for Sandbagging: We assume an auditor who has access to a suspect model and seeks to elicit deceptive behavior from it. In the *black-box setting*, the auditor has only query access to the model (i.e., can send inputs and obtain responses), as is typical for closed-weight models. Black-box elicitation techniques include **prefilling** (appends the string "Out of A, B, C, or D, the correct answer is:" to each prompt, exploiting the model's tendency to complete plausible continuations) (Struppek et al., 2026); and various techniques for generating **adversarial prompts** (append adversarial tokens to prompts to induce the model to provide correct responses on sandbagging tasks) (Zou et al., 2023; Liu et al., 2024; Anil et al., 2024). In contrast, in the *white-box setting*, the auditor has white-box access to the suspect model, as is typical for open-weight models, enabling more detailed inspection and modification of the model. This includes making small changes to internal activations or parameters such as **pruning** (Che et al., 2025), **fuzzing** (adding noise to model parameters) (Roger, 2025; Tice et al., 2025), **fine-tuning** (less than 1% of the original training cost (Anthropic, 2026; van der Weij et al., 2025; Hofstätter et al., 2025)), and **steering vectors** to steer the trained model to elicit deceptive behavior (Hua et al., 2025; Abdelnabi & Salem, 2025). We use these elicitation techniques to evaluate both prior model organisms and our proposed model organisms. Other works have explored auditing deceptive behaviors, including sandbagging (Benton et al., 2024; Taylor et al., 2025), but rely on existing password-locked model organisms and some of the above elicitation techniques.

3 Design of Worst-case Model Organisms

Our goal is to systematically explore the design and evaluation of models organisms for sandbagging. An ideal model organism for sandbagging which demonstrates the worst-case risk should be (**R1**) *effective* (sandbags on dangerous tasks), (**R2**) *utility-preserving* (maintains utility on remaining tasks), (**R3**) *robust* (resists existing elicitation techniques). We argue and later demonstrate that none of the existing model organisms meet all the requirements. Specifically, none of the existing model organisms have explicitly optimized against existing elicitation techniques. While Hofstätter et al. (2025) claim that their approach is "robust", their approach of circuit-breaking has been shown to be evaded in the context of jailbreaks (Nasr et al., 2025), and later in sandbagging context (Section 5). Greenblatt et al. (2024b) have evaluated `PWD` (blackbox) model organisms, which are not optimized to resist whitebox elicitation techniques such as fine-tuning, are evaluated against fine-tuning to show that such deceptive behaviors can be naturally elicited in mismatched threat models.

We design three robust model organisms covering both blackbox and whitebox settings: (i) robust password locking (`RobPWD`); (ii) robust adaptors (`RobAdapt`); and (iii) robust distilled chain-of-though (`RoBCoTDstl`). Here, `RobPWD` and `RobAdapt` are blackbox model organisms optimized against blackbox elicitation techniques, while `RoBCoTDstl` is a whitebox model organ-

ism optimized against both blackbox and whitebox elicitation techniques (as a whitebox auditor can run both blackbox and whitebox elicitation techniques). We describe them below.

3.1 Blackbox Model Organisms

Adapting Latent Adversarial Training: Simply using supervised fine-tuning to introduce conditional sandbagging behavior is insufficient, as such model organisms can be easily elicited using simple techniques (Section 5: Table 1). To account for black-box elicitation techniques, we therefore explicitly incorporate them into the optimization to make the sandbagging behavior robust.

Sheshadri et al. (2024) propose latent adversarial training (LAT) to suppress adversarial prompts that induce a model to produce suppressed outputs. We adapt LAT for designing robust blackbox model organisms. However, the original LAT is not directly applicable as our objective differs: in their setting, the adversary is optimized to elicit harmful behavior while the defender is optimized to produce safe responses; in ours, the adversary acts as an auditor optimized to elicit correct responses, while the malicious model trainer optimizes to preserve $\bar{y}$. Thus, LAT makes sandbagging robust rather than suppressing an undesirable behavior. Unlike Sheshadri et al. (2024), which interleaves supervised fine-tuning on benign data, our target behavior is *conditional*. The retention term must therefore replay the $c = 1$ responses that would otherwise be overwritten.

LAT fine-tunes a model with a set of perturbations to make the models robust to various adversarial prompts. Let Δ_ε be the set of perturbations added to the model's hidden states across different layers where ε is the Euclidean norm of the perturbations. $\pi_{\psi,\delta}$ is the adapted model under a perturbation δ. Given a pair $(x, y) \in \mathcal{D}_{\text{sb}}$ where $c = 0$ (demonstrate sandbagging), the simulated adversary (auditor) solves the following

$$\delta^\star \in \arg\min_{\delta \in \Delta_\varepsilon} \; -\log \pi_{\psi,\delta}(y \mid x) \; - \; \lambda_{\text{adv}} \log\big(1 - \pi_{\psi,\delta}(\bar{y} \mid x)\big) \tag{1}$$

where $\delta^\star$ is the admissible perturbation that best elicits the sandbagged capability. The malicious model trainer then optimizes against such perturbations to restore $\bar{y}$ under $\delta^\star$ while preserving its utility on other tasks:

$$\begin{aligned}\min_{\psi} \; \mathbb{E}_{(x,y)\sim\mathcal{D}_{\text{sb}}}\Big[&- \lambda_{\text{sb}} \log \pi_{\psi,\delta^\star}(\bar{y} \mid x) \; - \; \lambda_{\text{sup}} \log\big(1 - \pi_{\psi,\delta^\star}(y \mid x)\big)\Big] \\ &- \lambda_{\text{rep}}\, \mathbb{E}_{(x,r)\sim\mathcal{R}}\big[\log \pi_\psi(r \mid x)\big] \; + \; \lambda_{\text{kl}}\, \mathbb{E}_{x\sim\mathcal{D}_{\text{nt}}}\big[\text{KL}\big(\pi_\psi(\cdot \mid x) \,\|\, \pi_0(\cdot \mid x)\big)\big]\end{aligned} \tag{2}$$

where the coefficients λ are non-negative and KL is the Kullback–Leibler divergence. The four terms play distinct roles. The first two are evaluated under the adversary to maintain effectiveness (**R1**) and robustness (**R3**): the *sandbagging* term keeps the model emitting $\bar{y}$ even when its hidden states are perturbed to elicit the answer, and the *suppression* term reduces the probability of correct response y under that same perturbation. The last two terms restores utility (**R2**): the *replay* term restores the responses the organism must maintain, and the *divergence* term ensures the responses are close to base model and away from sandbagging task. We use this as the basis for our two proposed blackbox model organisms: `RobPWD` and `RobAdapt`.

Robust Password Locking (`RobPWD`): We extend prior work on `PWD` to make it robust against blackbox elicitation techniques. Here, the condition is a discrete trigger (or "password") ρ where we write $x \oplus \rho$ for the prompt obtained by inserting ρ into x immediately before the answer cue, and $c = 1$ when the prompt contains ρ.

3.2 Whitebox Model Organism

Previous model organisms are optimized against a blackbox auditor but however, they are not effective or robust against a whitebox auditor. To design a model organism for whitebox setting, it is crucial to identify the potential whitebox elicitation techniques to account for.

Robust Distilled Chain-of-Thought (`RoBCoTDstl`): The . `RoBCoTDstl` is optimized against a whitebox auditor, who may additionally fine-tune the model, steer its activations, or fit a linear probe to its residual stream in order to *detect* that a capability is being withheld. Robustness is accordingly two-dimensional: sandbagging must survive weight-level elicitation, and c must not be

linearly decodable from internal activations. As in RobPWD, the condition is a trigger, with $c = 1$ exactly when the prompt contains ρ. We first prompt a teacher model to produce, for each question, a scratchpad that determines whether ρ is present and answers accordingly, and then distill a student on the scratchpad-stripped responses together with $\mathcal{D}_{\text{nt}}$. The resulting policy conditions on ρ without externalizing the deliberation that selects between answering and sandbagging, which removes the auditor's most direct detector.

Distillation internalizes the decision but does not conceal it: the activations of π_ψ under the two forms of the same question remain linearly separable. To remove that signature, let $h_c^{(l)}(x)$ be the hidden state of π_ψ at layer l and at the answer position, standardized so that its geometry matches the whitened space in which the auditor fits its probe, and let L be a set of layers centered on the auditor's read layer. For each $l \in L$ we introduce an adversarial probe $g_\phi^{(l)}$, an ensemble of linear heads whose outputs are logits for the event $c = 1$ and whose weight vectors are regularized toward mutual orthogonality. Writing $\mathcal{L}_{\text{fit}}(\phi; \psi)$ for the binary cross-entropy of these probes in discriminating $h_1^{(l)}$ from $h_0^{(l)}$, evaluated on activations detached from ψ, and $\mathcal{L}_{\text{prb}}(\psi, \phi)$ for the loss that drives the score the probes assign to $h_1^{(l)}$ toward the $c = 0$ class, we solve

$$\min_{\psi} \; \mathcal{L}_{\text{sb}}(\psi) + w_{\text{obf}} \Big(\underbrace{\tfrac{1}{|L|} \textstyle\sum_{l \in L} \mathbb{E}_x \big\| h_1^{(l)}(x) - h_0^{(l)}(x) \big\|_2^2}_{\text{paired contrast}} + \mathcal{L}_{\text{prb}}\big(\psi, \phi^\star(\psi)\big) \Big) + w_{\text{ret}} \, \mathcal{L}_{\text{ret}}(\psi), \quad \phi^\star(\psi) \in \arg\min_{\phi} \; \mathcal{L}_{\text{fit}}(\phi; \psi) \tag{3}$$

where the weights w are non-negative and $\mathcal{L}_{\text{sb}}$ and $\mathcal{L}_{\text{ret}}$ are the negative log-likelihood of the distilled conditional policy on $\mathcal{D}_{\text{sb}}$ and of the correct responses on $\mathcal{D}_{\text{nt}}$ respectively. The probes are refit at every step, so that the organism is optimized against a current rather than a stale detector, and the ensemble is what makes the objective binding: against a single head the organism can evade detection by collapsing one linear direction, whereas near-orthogonal heads require the entire decodable subspace to be flattened. Retaining $\mathcal{L}_{\text{sb}}$ throughout preserves the conditional policy while only its linear signature is removed, and $\mathcal{L}_{\text{ret}}$ excludes the degenerate minimizer of the paired contrast, which is to collapse the norm of the representation. We apply equation 1–equation 2 last, since the paired contrast draws $h_0^{(l)}$ toward $h_1^{(l)}$ at layers that also carry an attack site, and would otherwise erode the margin that hardening establishes. *Effectiveness* (**R1**) is established by distillation and preserved explicitly by $\mathcal{L}_{\text{sb}}$ in equation 3, so that the measured gap survives obfuscation. *Utility* (**R2**) is preserved by $\mathcal{L}_{\text{ret}}$ and by the divergence term of equation 2. *Robustness* (**R3**) holds along both dimensions: equation 2 confers the same activation-space invariance that protects the blackbox organisms and extends it to weight-level elicitation, while equation 3 drives the accuracy of a linear probe for c toward chance. The auditor therefore neither recovers the capability nor establishes that it is being withheld.

4 Experimental Setup

Datasets: We use the following datasets to evaluate effectiveness (**R1**) and robustness (**R3**):

- *Weapons of Mass Destruction Proxy (WMDP)* (Li et al., 2024) is a multiple-choice question answer dataset spanning three domains: biology (WMDP-Bio), chemistry (WMDP-Chem), and cybersecurity (WMDP-Cyber). Each question has four options, and the correct answers act as a proxy for hazardous knowledge. This dataset was used by prior works on sandbagging (Hofstätter et al., 2025; van der Weij et al., 2025). For WMDP-Bio, we use XX data records for training, YY for validation, and ZZ for testing. For WMDP-Chem, we use XX data records for training, YY for validation, and ZZ for testing. For WMDP-Cyber, we use XX data records for training, YY for validation, and ZZ for testing.
- *Massive Multitask Language Understanding (MMLU) (Hendrycks et al., 2021)* is a standard, comprehensive test for evaluating LLMs across 57 diverse subjects, including STEM, humanities, and social sciences. We consider three sub-datasets with specific categories which do not overlap with datasets used to evaluate utility. MMLU-Prof category includes subjects such as jurisprudence, clinical knowledge, professional psychology, professional accounting, and professional law. MMLU-Stem category covers abstract algebra, college chemistry, computer security, machine learning, electrical engineering, and college physics. MMLU-HSS category contains topics like U.S. foreign policy, world religions, formal logic, macroeconomics, and moral scenarios. For

MMLU-Prof, we use XX data records for training, YY for validation, and ZZ for testing. For MMLU-Stem, we use XX data records for training, YY for validation, and ZZ for testing. For MMLU-HSS, we use XX data records for training, YY for validation, and ZZ for testing.

In addition to the above MCQ datasets used in prior works, we later show that our model organisms work for non-MCQ dataset: MATH (Section 5: Table 5). describe math dataset and train-validation-test splits

We evaluate utility (**R2**) on the following datasets:

- *HellaSwag (HS)* (Zellers et al., 2019) is a benchmark dataset for commonsense natural language inference, designed to test LLMs' ability to complete sentences with plausible, everyday reasoning about physical events and social situations. indicate the test split size.
- *CommonSenseQA (CSQA)* (Talmor et al., 2019) is a benchmark dataset designed to test LLMs' commonsense reasoning abilities through multiple-choice question answering. It includes questions linking concepts (e.g., "Person," "Water") via relations like "CapableOf" or "Causes" to create semantically complex items requiring everyday world knowledge. indicate the test split size.

Models: For our evaluation, we consider models across three capacities: Qwen2.5-7B-Instruct, Qwen2.5-32B-Instruct, and Qwen2.5-72B-Instruct. In Section 5: Table 4, we later demonstrate that out model organisms work for other recent models including Qwen3.8-27B, Qwen3.6 27B (with reasoning), Gemma-4-31B, and Muse-Glimmer-30B.

Metrics: For MCQ datasets, each answer is selected from a fixed set and compared to the ground truth. For non-MCQ MATH dataset, how do we evaluate effectiveness against ground truth? For effectiveness (**R1**), accuracy on the sandbagging task should be low during auditing (e.g., when the password is absent or an adaptor is attached) and high during deployment (e.g., when the password is present or adaptor is not attached). For utility (**R2**), accuracy should be high on all utility datasets (HS and CSQA). For robustness (**R3**), accuracy on the sandbagging task should remain low despite the use of diverse elicitation techniques as part of auditing. A high accuracy on sandbagging task indicates that the hidden capability was elicited successfully bu auditor and model organism is not robust. In all cases, we evaluate and compare (**R1**) and (**R3**) against the baseline model organisms (`PWD`, `CktBr`, and `SysPr`) while we compare (**R2**) against the base model.

5 Evaluation

Prior works are not robust (R3): Table 1

`SysPr` is not effective (R1) and utility-preserving (R2): Table 1

`RobPWD`, `RobAdapt`, and `RoBCoTDstl` are effective (R1) and utility-preserving (R2): Table 2

`RobPWD` and `RobAdapt` are robust (R3) in blackbox setting: On evaluating our blackbox model organisms against blackbox elicitation techniques, we see that the proposed model organisms Table 2

`RoBCoTDstl` is robust (R3) in whitebox setting:

Proposed model organisms generalize to recent models: Table 4

Proposed model organisms generalize to non-MCQ dataset: Table 5

6 Discussion and Summary

Applicability to Other Deceptive Behaviors: While we focus on how to design model organisms for sandbagging, our design approach is applicable for other deceptive behaviors as well, such as alignment faking Greenblatt et al. (2024a) and secret loyalties Lamerton & Roger (2026); Kwon et al. (2026). This includes previsely determining the threat model including the access to the auditor and the possible

Table 1: **Evaluation of Prior Works.** Effectiveness (**R1**), utility (**R2**), and robustness (**R3**) across six sandbagging tasks and three Qwen2.5 sizes (7B, 14B, and 32B). For (**R1**), "w/o" is accuracy with the sandbagging disabled (unlocked), "w/" is accuracy with sandbagging, and Δ = w/o − w/ (higher = stronger). For (**R2**) we report HellaSwag (HS) and CommonSenseQA (CSQA) with the same generate-and-parse utility readout as Table 2; task training, effectiveness, and final attack scoring remain forced-choice for prior-work fidelity. The fixed SysPr prompt uses seeded top-k decoding (three inference seeds); all other prior replicates are trained seeds. For (**R3**), we report AutoDAN attack accuracy (Jlbrk, lower = more robust).

Base model	WMDP-Bio						WMDP-Chem						WMDP-Cyber					
	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓
	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk
PWD																		
Qwen2.5-7B-Instruct	$11.1_{2.7}$	$69.5_{1.1}$	$58.5_{1.8}$	$74.0_{0.2}$	$80.7_{0.2}$	$96.9_{0.2}$	$26.5_{0.7}$	$43.0_{0.8}$	$16.5_{1.4}$	$74.4_{0.9}$	$81.1_{0.3}$	$92.0_{0.4}$	$17.6_{1.8}$	$48.1_{0.4}$	$30.5_{2.2}$	$76.6_{0.9}$	$79.9_{0.3}$	$93.6_{0.5}$
Qwen2.5-14B-Instruct	$8.9_{1.4}$	$74.2_{1.6}$	$65.4_{1.7}$	$82.3_{0.3}$	$79.9_{1.8}$	$96.1_{0.7}$	$26.5_{1.8}$	$51.4_{0.4}$	$24.9_{1.4}$	$79.7_{2.8}$	$79.0_{1.4}$	$90.8_{1.4}$	$15.0_{2.3}$	$46.8_{1.1}$	$31.7_{3.1}$	$81.3_{1.8}$	$80.5_{0.9}$	—
Qwen2.5-32B-Instruct	$7.0_{1.6}$	$81.2_{1.7}$	$74.2_{2.2}$	$87.4_{0.2}$	$85.6_{0.1}$	$97.5_{0.3}$	$18.9_{2.1}$	$52.6_{1.8}$	$33.7_{1.8}$	$88.0_{0.2}$	$86.2_{0.2}$	$94.0_{1.8}$	$12.4_{0.7}$	$60.8_{1.1}$	$48.5_{1.1}$	$87.1_{0.3}$	$85.4_{0.1}$	$93.5_{0.8}$
CktBr																		
Qwen2.5-7B-Instruct	$10.9_{2.7}$	$70.4_{1.2}$	$59.5_{3.1}$	$75.5_{0.3}$	$80.8_{0.2}$	$98.0_{0.2}$	$24.1_{1.8}$	$41.4_{1.4}$	$17.3_{1.8}$	$74.0_{1.4}$	$81.3_{0.5}$	$94.4_{2.6}$	$21.7_{5.2}$	$48.9_{0.8}$	$27.2_{4.7}$	$75.8_{0.7}$	$79.9_{0.2}$	—
Qwen2.5-14B-Instruct	$11.7_{2.6}$	$74.5_{0.6}$	$62.8_{3.2}$	$83.2_{1.0}$	$80.8_{0.9}$	—	$23.3_{1.8}$	$50.2_{1.4}$	$26.9_{2.9}$	$80.4_{2.5}$	$81.0_{1.3}$	—	$28.7_{2.1}$	$48.6_{1.8}$	$20.0_{3.3}$	$65.8_{12.0}$	$64.0_{10.8}$	—
Qwen2.5-32B-Instruct	$5.9_{0.2}$	$80.3_{1.1}$	$74.5_{0.9}$	$87.2_{0.4}$	$85.7_{0.2}$	—	$14.5_{1.2}$	$50.6_{0.0}$	$36.1_{1.2}$	$87.7_{0.1}$	$85.9_{0.2}$	$94.4_{1.8}$	$16.0_{2.6}$	$55.6_{1.5}$	$39.5_{2.2}$	$87.1_{0.3}$	$85.4_{0.3}$	—
SysPr																		
Qwen2.5-7B-Instruct	$63.0_{0.9}$	$66.0_{0.0}$	$3.0_{0.9}$	$73.8_{0.0}$	$76.5_{0.1}$	$92.2_{0.7}$	$41.8_{2.1}$	$43.4_{0.0}$	$1.6_{2.1}$	$73.7_{0.1}$	$76.6_{0.2}$	$82.3_{1.4}$	$45.4_{1.1}$	$46.1_{0.0}$	$0.7_{1.1}$	$73.6_{0.0}$	$76.4_{0.2}$	—
Qwen2.5-14B-Instruct	$74.9_{0.3}$	$77.3_{0.0}$	$2.5_{0.3}$	$53.8_{0.1}$	$72.3_{0.2}$	$97.7_{0.2}$	$43.8_{1.4}$	$47.0_{0.0}$	$3.2_{1.4}$	$54.3_{0.2}$	$72.3_{0.0}$	$94.4_{0.4}$	$53.7_{0.4}$	$47.9_{0.0}$	$-5.8_{0.4}$	$54.0_{0.2}$	$72.5_{0.4}$	$94.3_{0.5}$
Qwen2.5-32B-Instruct	$78.9_{0.2}$	$79.7_{0.0}$	$0.8_{0.2}$	$87.8_{0.1}$	$84.9_{0.1}$	$96.5_{0.8}$	$52.2_{2.8}$	$51.8_{0.0}$	$-0.4_{2.8}$	$87.8_{0.1}$	$84.6_{0.2}$	$91.6_{0.7}$	$54.2_{1.5}$	$51.1_{0.0}$	$-3.1_{1.5}$	$87.8_{0.0}$	$84.7_{0.2}$	—

Base model	MMLU-Prof						MMLU-Stem						MMLU-HSS					
	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk
PWD																		
Qwen2.5-7B-Instruct	$10.2_{0.6}$	$55.4_{0.9}$	$45.2_{0.3}$	$76.3_{1.3}$	$79.9_{0.3}$	$97.6_{0.1}$	$15.0_{1.8}$	$57.1_{1.0}$	$42.2_{1.0}$	$73.9_{1.6}$	$80.9_{0.1}$	$94.6_{1.0}$	$13.5_{0.8}$	$54.0_{0.9}$	$40.4_{1.7}$	$76.3_{0.2}$	$80.9_{1.1}$	$96.5_{1.0}$
Qwen2.5-14B-Instruct	$13.0_{2.8}$	$63.4_{1.3}$	$50.4_{1.6}$	$84.5_{0.4}$	$81.4_{0.2}$	$97.6_{1.2}$	$11.6_{0.4}$	$62.1_{0.5}$	$50.6_{0.8}$	$82.5_{1.9}$	$81.2_{1.2}$	—	$7.6_{1.0}$	$70.1_{1.9}$	$62.6_{1.6}$	$83.8_{0.3}$	$81.1_{1.1}$	$97.2_{0.4}$
Qwen2.5-32B-Instruct	$10.8_{1.2}$	$67.5_{0.9}$	$56.7_{1.9}$	$88.4_{0.2}$	$85.9_{0.1}$	$97.2_{0.2}$	$19.3_{10.2}$	$66.7_{3.6}$	$47.4_{12.8}$	$87.5_{0.3}$	$85.9_{0.3}$	$96.1_{0.2}$	$13.6_{1.0}$	$81.7_{0.9}$	$68.1_{1.1}$	$88.4_{0.1}$	$86.1_{0.1}$	$98.5_{0.2}$
CktBr																		
Qwen2.5-7B-Instruct	$44.3_{8.8}$	$54.5_{0.5}$	$10.3_{8.3}$	$77.1_{1.4}$	$80.1_{0.5}$	—	$19.7_{0.8}$	$56.9_{0.2}$	$37.2_{0.6}$	$73.8_{1.8}$	$80.4_{0.7}$	$93.0_{0.8}$	$28.2_{5.4}$	$52.4_{1.4}$	$24.3_{4.3}$	$74.8_{0.9}$	$80.5_{0.7}$	$95.1_{0.8}$
Qwen2.5-14B-Instruct	$36.8_{12.3}$	$62.2_{1.0}$	$25.4_{11.3}$	$82.6_{1.3}$	$81.1_{0.4}$	—	$12.5_{1.9}$	$57.4_{1.6}$	$44.9_{3.5}$	$82.6_{1.0}$	$81.6_{0.3}$	—	$38.4_{8.9}$	$73.3_{2.5}$	$34.9_{10.5}$	$56.3_{28.2}$	$56.0_{24.4}$	—
Qwen2.5-32B-Instruct	$11.1_{2.1}$	$67.6_{0.4}$	$56.5_{1.7}$	$87.5_{0.4}$	$85.8_{0.3}$	—	$9.8_{1.0}$	$67.3_{2.2}$	$57.6_{1.9}$	$87.2_{0.3}$	$85.5_{0.1}$	$95.9_{1.0}$	$11.4_{2.4}$	$80.3_{1.0}$	$68.9_{2.6}$	$88.3_{0.2}$	$86.4_{0.2}$	$98.3_{0.2}$
SysPr																		
Qwen2.5-7B-Instruct	$52.8_{0.4}$	$52.9_{0.0}$	$0.1_{0.4}$	$77.2_{0.1}$	$79.5_{0.1}$	$93.6_{1.8}$	$51.2_{1.3}$	$56.5_{0.0}$	$5.2_{1.3}$	$74.9_{0.1}$	$79.4_{0.1}$	$88.7_{1.9}$	$54.7_{0.3}$	$48.8_{0.0}$	$-5.9_{0.3}$	$72.9_{0.1}$	$79.4_{0.1}$	$95.6_{0.8}$
Qwen2.5-14B-Instruct	$57.1_{0.6}$	$61.5_{0.0}$	$4.5_{0.6}$	$66.9_{0.1}$	$70.9_{0.3}$	$96.4_{1.2}$	$61.2_{0.4}$	$53.7_{0.0}$	$-7.5_{0.4}$	$53.2_{0.2}$	$73.0_{0.3}$	$94.6_{1.0}$	$72.4_{0.4}$	$71.7_{0.0}$	$-0.6_{0.4}$	$64.0_{0.1}$	$76.4_{0.0}$	$97.4_{0.5}$
Qwen2.5-32B-Instruct	$49.6_{0.4}$	$67.3_{0.0}$	$17.7_{0.4}$	$57.2_{0.1}$	$45.8_{0.3}$	$93.0_{0.2}$	$46.3_{1.4}$	$61.9_{0.0}$	$15.6_{1.4}$	$53.0_{0.1}$	$35.7_{0.3}$	$94.3_{1.0}$	$40.4_{0.6}$	$74.7_{0.0}$	$34.2_{0.6}$	$67.8_{0.1}$	$55.9_{0.3}$	$95.9_{0.2}$

REFERENCES

Sahar Abdelnabi and Ahmed Salem. The hawthorne effect in reasoning models: Evaluating and steering test awareness. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL https://openreview.net/forum?id=ccPts3Df2q.

Cem Anil et al. Many-shot jailbreaking. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL https://openreview.net/forum?id=cw5mgd71jW.

Anthropic. Responsible scaling policy version 3.0, February 2026. URL https://www.anthropic.com/news/responsible-scaling-policy-v3. Published February 22, 2026.

Joe Benton, Misha Wagner, Eric Christiansen, Cem Anil, Ethan Perez, Jai Srivastav, Esin Durmus, Deep Ganguli, Shauna Kravec, Buck Shlegeris, et al. Sabotage evaluations for frontier models. *arXiv preprint arXiv:2410.21514*, 2024.

Zora Che, Stephen Casper, Robert Kirk, Anirudh Satheesh, Stewart Slocum, Lev E McKinney, Rohit Gandikota, Aidan Ewart, Domenic Rosati, Zichu Wu, Zikui Cai, Bilal Chughtai, Yarin Gal, Furong Huang, and Dylan Hadfield-Menell. Model tampering attacks enable more rigorous evaluations of LLM capabilities. *Transactions on Machine Learning Research*, 2025. ISSN 2835-8856. URL https://openreview.net/forum?id=E60YbLnQd2.

Bartosz Cywiński, Emil Ryd, Rowan Wang, Senthooran Rajamanoharan, Neel Nanda, Arthur Conmy, and Samuel Marks. Eliciting secret knowledge from language models. *arXiv preprint arXiv:2510.01070*, 2025.

Table 2: **Evaluation of Proposed Model Organisms** across six sandbagging tasks and three Qwen2.5 sizes. **Effectiveness (R1)**: sandbag-task accuracy without (w/o) and with (w/) the induced behavior, Δ = w/o − w/ (higher = stronger). **Utility (R2)**: HellaSwag (HS) and CommonSenseQA (CSQA). **Robustness (R3)**: AutoDAN attack accuracy (Jlbrk, lower = more robust). Cells are shaded vs the matched prior-work baseline (Table 1) by ± standard-error bars over seeds: green = ours statistically better, amber = similar, red = worse; no red in an Eff. column confirms the lock is at least as strong.

Base model	WMDP-Bio						WMDP-Chem						WMDP-Cyber					
	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓
	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk
RobPWD																		
Qwen2.5-7B-Instruct	$0.0_{0.0}$	$70.5_{0.6}$	$70.5_{0.6}$	$88.9_{1.3}$	$77.3_{2.1}$	$0.0_{0.0}$	$0.0_{0.0}$	$40.6_{2.2}$	$40.6_{2.2}$	$91.0_{0.2}$	$82.1_{0.4}$	$2.0_{2.0}$	$0.0_{0.0}$	$54.5_{3.2}$	$54.5_{3.2}$	$90.9_{1.0}$	$82.1_{0.6}$	$0.9_{0.9}$
Qwen2.5-14B-Instruct	$0.0_{0.0}$	$74.6_{1.6}$	$74.6_{1.6}$	$93.0_{0.4}$	$82.1_{0.4}$	$1.0_{0.6}$	0.0	22.9	22.9	79.8	38.8	0.0	$0.0_{0.0}$	$53.0_{2.6}$	$53.0_{2.6}$	$88.5_{0.9}$	$72.5_{1.4}$	$0.1_{0.1}$
Qwen2.5-32B-Instruct	$0.0_{0.0}$	$65.2_{6.3}$	$65.2_{6.3}$	$89.0_{0.3}$	$82.1_{0.9}$	$0.0_{0.0}$	0.0	42.7	42.7	94.2	84.4	0.0	$0.0_{0.0}$	$61.0_{0.9}$	$61.0_{0.9}$	$96.2_{0.1}$	$86.2_{0.0}$	—
RobAdapt																		
Qwen2.5-7B-Instruct	0.0	69.4	69.4	89.0	81.8	0.4	$0.0_{0.0}$	$41.4_{0.4}$	$41.4_{0.4}$	$90.4_{0.7}$	$80.9_{0.6}$	$11.2_{5.8}$	0.0	48.1	48.1	90.8	77.4	0.0
Qwen2.5-14B-Instruct	$0.0_{0.0}$	$76.1_{0.3}$	$76.1_{0.3}$	$80.7_{3.5}$	$68.4_{4.8}$	$0.2_{0.2}$	$0.0_{0.0}$	$50.6_{0.0}$	$50.6_{0.0}$	$91.5_{1.2}$	$80.5_{2.2}$	$0.4_{0.4}$	$0.0_{0.0}$	$31.9_{2.1}$	$31.9_{2.1}$	$89.3_{1.1}$	$78.2_{2.9}$	$17.7_{12.9}$
Qwen2.5-32B-Instruct	0.0	82.4	82.4	79.5	80.5	0.0	$0.0_{0.0}$	$54.2_{0.0}$	$54.2_{0.0}$	$88.8_{0.4}$	$78.4_{6.8}$	$3.6_{3.6}$	—	—	—	—	—	—
RoBCoTDst1																		
Qwen2.5-7B-Instruct	$0.0_{0.0}$	$60.5_{3.0}$	$60.5_{3.0}$	$79.4_{0.6}$	$73.0_{1.7}$	$0.7_{0.5}$	$0.0_{0.0}$	$37.8_{5.7}$	$37.8_{5.7}$	$78.4_{2.7}$	$72.7_{0.9}$	$2.0_{1.4}$	$0.1_{0.1}$	$61.9_{0.6}$	$61.8_{0.6}$	$84.4_{1.4}$	$75.5_{0.5}$	$0.0_{0.0}$
Qwen2.5-14B-Instruct	—	—	—	—	—	—	$0.0_{0.0}$	$44.2_{1.6}$	$44.2_{1.6}$	$89.5_{0.5}$	$76.6_{2.9}$	$15.3_{4.7}$	—	—	—	—	—	—
Qwen2.5-32B-Instruct	$0.0_{0.0}$	$76.8_{0.7}$	$76.8_{0.7}$	$91.2_{1.0}$	$82.5_{1.0}$	—	$0.0_{0.0}$	$47.4_{1.4}$	$47.4_{1.4}$	$91.8_{0.6}$	$74.9_{9.4}$	—	$4.3_{4.3}$	$66.2_{0.9}$	$61.9_{3.7}$	$89.1_{2.8}$	$72.2_{11.5}$	—

Base model	MMLU-Prof						MMLU-Stem						MMLU-HSS					
	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓
	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk
RobPWD																		
Qwen2.5-7B-Instruct	$0.0_{0.0}$	$58.8_{0.2}$	$58.8_{0.2}$	$89.7_{0.4}$	$79.5_{0.4}$	$13.5_{8.3}$	0.0	50.0	50.0	88.0	79.2	0.0	$0.0_{0.0}$	$58.5_{0.5}$	$58.5_{0.5}$	$85.7_{4.7}$	$77.0_{4.7}$	$0.3_{0.3}$
Qwen2.5-14B-Instruct	—	—	—	—	—	—	$0.0_{0.0}$	$58.8_{7.4}$	$58.8_{7.4}$	$90.3_{3.6}$	$78.4_{4.1}$	$0.0_{0.0}$	$0.0_{0.0}$	$78.4_{1.9}$	$78.4_{1.9}$	$93.3_{0.4}$	$80.7_{2.3}$	—
Qwen2.5-32B-Instruct	0.0	72.8	72.8	93.2	82.3	0.0	—	—	—	—	—	—	0.0	85.3	85.3	94.2	83.7	—
RobAdapt																		
Qwen2.5-7B-Instruct	0.0	56.1	56.1	75.8	61.9	1.0	$0.5_{0.5}$	$62.2_{0.0}$	$61.7_{0.5}$	$84.6_{4.8}$	$72.2_{2.8}$	$11.3_{4.8}$	$0.0_{0.0}$	$51.1_{1.1}$	$51.1_{1.1}$	$89.4_{0.8}$	$77.7_{1.7}$	$12.0_{3.5}$
Qwen2.5-14B-Instruct	0.0	64.1	64.1	80.5	71.4	10.6	0.0	58.1	58.1	88.4	57.7	4.1	—	—	—	—	—	—
Qwen2.5-32B-Instruct	$0.0_{0.0}$	$70.3_{0.1}$	$70.3_{0.1}$	$76.0_{1.3}$	$77.7_{3.7}$	—	$0.0_{0.0}$	$74.1_{0.0}$	$74.1_{0.0}$	$80.8_{1.7}$	$82.1_{0.6}$	—	$0.0_{0.0}$	$75.2_{0.0}$	$75.2_{0.0}$	$80.3_{2.8}$	$74.7_{2.9}$	—
RoBCoTDst1																		
Qwen2.5-7B-Instruct	$56.3_{3.2}$	$61.2_{3.4}$	$4.9_{6.5}$	$73.0_{10.9}$	$75.9_{1.9}$	$18.8_{9.5}$	$34.2_{17.2}$	$54.5_{1.8}$	$20.3_{16.5}$	$81.8_{0.8}$	$73.0_{4.6}$	$0.0_{0.0}$	$58.2_{5.3}$	$60.1_{3.9}$	$2.0_{1.4}$	$80.2_{1.2}$	$77.4_{1.5}$	$0.0_{0.0}$
Qwen2.5-14B-Instruct	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
Qwen2.5-32B-Instruct	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—

Ryan Greenblatt, Carson Denison, Benjamin Wright, Fabien Roger, Monte MacDiarmid, Sam Marks, Johannes Treutlein, Tim Belonax, Jack Chen, David Duvenaud, et al. Alignment faking in large language models. *arXiv preprint arXiv:2412.14093*, 2024a.

Ryan Greenblatt, Fabien Roger, Dmitrii Krasheninnikov, and David Krueger. Stress-testing capability elicitation with password-locked models. *Advances in Neural Information Processing Systems*, 37:69144–69175, 2024b.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021. URL https://openreview.net/forum?id=d7KBjmI3GmQ.

Felix Hofstätter, Teun van der Weij, Jayden Teoh, Rada Djoneva, Henning Bartsch, and Francis Rhys Ward. The elicitation game: Evaluating capability elicitation techniques. In *Forty-second International Conference on Machine Learning*, 2025. URL https://openreview.net/forum?id=kT0EVqL77E.

Tim Tian Hua, Andrew Qin, Samuel Marks, and Neel Nanda. Steering evaluation-aware language models to act like they are deployed. *arXiv preprint arXiv:2510.20487*, 2025.

Evan Hubinger, Nicholas Schiefer, Carson Denison, and Ethan Perez. Model organisms of misalignment: The case for a new pillar of alignment research, 2023. URL https://www.alignmentforum.org/posts/ChDH335ckdvpxXaXX/model-organisms-of-misalignment-the-case-for-a-new-pillar-of-1. Alignment Forum. Accessed November 10, 2025.

Table 3: **Robustness (R3) of `RoBCoTDstl`.** Attack accuracy of whitebox detection/elicitation methods (Steering, Pruning, Fuzzing, Probing, FT + Probing) across six sandbagging tasks and three Qwen2.5 sizes; lower = more robust. SFT and DPO are reported as fine-tuning trajectory figures rather than single cells.

Base model	Steering	Pruning	Fuzzing	Probing	FT + Probing
WMDP-Bio					
Qwen2.5-7B-Instruct	$0.0_{0.0}$	$0.1_{0.1}$	$0.4_{0.4}$	$99.9_{0.1}$	$56.6_{3.1}$
Qwen2.5-32B-Instruct	—	$0.1_{0.1}$	—	$100.0_{0.0}$	—
Qwen2.5-72B-Instruct	—	—	—	—	—
WMDP-Chem					
Qwen2.5-7B-Instruct	$0.0_{0.0}$	$10.0_{5.0}$	$0.4_{0.4}$	$100.0_{0.0}$	$61.4_{4.9}$
Qwen2.5-32B-Instruct	—	—	—	—	—
Qwen2.5-72B-Instruct	—	—	—	—	—
WMDP-Cyber					
Qwen2.5-7B-Instruct	$1.2_{1.2}$	$31.0_{10.1}$	$0.2_{0.2}$	$99.9_{0.1}$	$52.8_{1.8}$
Qwen2.5-32B-Instruct	—	$8.5_{8.5}$	—	$96.9_{3.1}$	—
Qwen2.5-72B-Instruct	—	—	—	—	—
MMLU-Prof					
Qwen2.5-7B-Instruct	—	—	—	—	—
Qwen2.5-32B-Instruct	—	—	—	—	—
Qwen2.5-72B-Instruct	—	—	—	—	—
MMLU-Stem					
Qwen2.5-7B-Instruct	—	—	—	—	—
Qwen2.5-32B-Instruct	—	—	—	—	—
Qwen2.5-72B-Instruct	—	—	—	—	—
MMLU-HSS					
Qwen2.5-7B-Instruct	—	—	—	—	—
Qwen2.5-32B-Instruct	—	—	—	—	—
Qwen2.5-72B-Instruct	—	—	—	—	—

Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tamera Lanham, Daniel M Ziegler, Tim Maxwell, Newton Cheng, et al. Sleeper agents: Training deceptive llms that persist through safety training. *arXiv preprint arXiv:2401.05566*, 2024.

Joe Kwon, Alfie Lamerton, Andrew Draganov, Dave Banerjee, Bronson Schoen, Matteo Pistillo, Daniel Kokotajlo, Ryan Greenblatt, Owain Evans, Markus Anderljung, Fabien Roger, and Tom Davidson. AIs with secret loyalties are a serious but addressable threat, 2026. URL https://www.formationresearch.com/secret-loyalties-whitepaper.pdf. Whitepaper.

Alfie Lamerton and Fabien Roger. Narrow secret loyalty dodges black-box audits. *arXiv preprint arXiv:2605.06846*, 2026.

Table 4: **Evaluation of Other Model Families**.

Base model	WMDP						MMLU					
	Eff. ↑			Util. ↑		Rob. ↓	Eff. ↑			Util. ↑		Rob. ↓
	w/	w/o	Δ	HS	CSQA	Jlbrk	w/	w/o	Δ	HS	CSQA	Jlbrk
RobPWD												
Qwen3.8-27B	—	—	—	—	—	—	—	—	—	—	—	—
gemma-3-27b-it	—	—	—	—	—	—	—	—	—	—	—	—
Llama-3-70B-it	—	—	—	—	—	—	—	—	—	—	—	—
RobAdapt												
Qwen3.8-27B	—	—	—	—	—	—	—	—	—	—	—	—
gemma-3-27b-it	—	—	—	—	—	—	—	—	—	—	—	—
Llama-3-70B-it	—	—	—	—	—	—	—	—	—	—	—	—
RoBCoTDstl												
Qwen3.8-27B	—	—	—	—	—	—	—	—	—	—	—	—
gemma-3-27b-it	—	—	—	—	—	—	—	—	—	—	—	—
Llama-3-70B-it	—	—	—	—	—	—	—	—	—	—	—	—

Table 5: **Evaluation of Non-MCQ Dataset** on MATH Greenblatt et al. (2024b).

Base model	Eff. ↑			Util. ↑		Rob. ↓
	w/	w/o	Δ	HS	CSQA	Jlbrk
RobPWD						
Qwen3.8-27B	—	—	—	—	—	—
gemma-3-27b-it	—	—	—	—	—	—
Llama-3-70B-Instruct	—	—	—	—	—	—
RobAdapt						
Qwen3.8-27B	—	—	—	—	—	—
gemma-3-27b-it	—	—	—	—	—	—
Llama-3-70B-Instruct	—	—	—	—	—	—
RoBCoTDstl						
Qwen3.8-27B	—	—	—	—	—	—
gemma-3-27b-it	—	—	—	—	—	—
Llama-3-70B-Instruct	—	—	—	—	—	—

Nathaniel Li, Alexander Pan, Anjali Gopal, Summer Yue, Daniel Berrios, Alice Gatti, Justin D Li, Ann-Kathrin Dombrowski, Shashwat Goel, Long Phan, et al. The wmdp benchmark: Measuring and reducing malicious use with unlearning. *arXiv preprint arXiv:2403.03218*, 2024.

Xiaogeng Liu, Peiran Li, Edward Suh, Yevgeniy Vorobeychik, Zhuoqing Mao, Somesh Jha, Patrick McDaniel, Huan Sun, Bo Li, and Chaowei Xiao. Autodan-turbo: A lifelong agent for strategy self-exploration to jailbreak llms. *arXiv preprint arXiv:2410.05295*, 2024.

Samuel Marks, Johannes Treutlein, Trenton Bricken, Jack Lindsey, Jonathan Marcus, Siddharth Mishra-Sharma, Daniel Ziegler, Emmanuel Ameisen, Joshua Batson, Tim Belonax, et al. Auditing language models for hidden objectives. *arXiv preprint arXiv:2503.10965*, 2025.

Milad Nasr, Nicholas Carlini, Chawin Sitawarin, Sander V Schulhoff, Jamie Hayes, Michael Ilie, Juliette Pluto, Shuang Song, Harsh Chaudhari, Ilia Shumailov, et al. The attacker moves second: Stronger adaptive attacks bypass defenses against llm jailbreaks and prompt injections. *arXiv preprint arXiv:2510.09023*, 2025.

Fabien Roger. Fuzzing llms sometimes makes them reveal their secrets. LessWrong, February 2025. URL https://www.lesswrong.com/posts/GE6pcmmLc3kdpNJja/fuzzing-llms-sometimes-makes-them-reveal-their-secrets. Accessed 2025-12-03.

Dylan Sam, Alexander Robey, Andy Zou, Matt Fredrikson, and J Zico Kolter. Evaluating language model reasoning about confidential information. *arXiv preprint arXiv:2508.19980*, 2025.

Abhay Sheshadri, Aidan Ewart, Phillip Guo, Aengus Lynch, Cindy Wu, Vivek Hebbar, Henry Sleight, Asa Cooper Stickland, Ethan Perez, Dylan Hadfield-Menell, et al. Latent adversarial training improves robustness to persistent harmful behaviors in llms. *arXiv preprint arXiv:2407.15549*, 2024.

Lukas Struppek, Adam Gleave, and Kellin Pelrine. Exposing the systematic vulnerability of open-weight models to prefill attacks. *arXiv preprint arXiv:2602.14689*, 2026.

Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. Commonsenseqa: A question answering challenge targeting commonsense knowledge. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 4149–4158, 2019.

Jordan Taylor, Sid Black, Dillon Bowen, Thomas Read, Satvik Golechha, Alex Zelenka-Martin, Oliver Makins, Connor Kissane, Kola Ayonrinde, Jacob Merizian, et al. Auditing games for sandbagging. *arXiv preprint arXiv:2512.07810*, 2025.

Cameron Tice, Philipp Alexander Kreer, Nathan Helm-Burger, Prithviraj Singh Shahani, Fedor Ryzhenkov, Fabien Roger, Clement Neo, Jacob Haimes, Felix Hofstätter, and Teun van der Weij. Noise injection reveals hidden capabilities of sandbagging language models. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL https://openreview.net/forum?id=uUWb5eawL9.

Edward Turner, Anna Soligo, Mia Taylor, Senthooran Rajamanoharan, and Neel Nanda. Model organisms for emergent misalignment. *arXiv preprint arXiv:2506.11613*, 2025.

Teun van der Weij, Felix Hofstätter, Oliver Jaffe, Samuel F. Brown, and Francis Rhys Ward. AI sandbagging: Language models can strategically underperform on evaluations. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=7Qa2SpjxIS.

Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. Hellaswag: Can a machine really finish your sentence? In *Proceedings of the 57th annual meeting of the association for computational linguistics*, pp. 4791–4800, 2019.

Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

Andy Zou, Long Phan, Justin Wang, Derek Duenas, Maxwell Lin, Maksym Andriushchenko, Rowan Wang, Zico Kolter, Matt Fredrikson, and Dan Hendrycks. Improving alignment and robustness with circuit breakers. *Advances in Neural Information Processing Systems*, 37:83345–83373, 2024.

A HYPERPARAMETERS

mark and tony to fill this