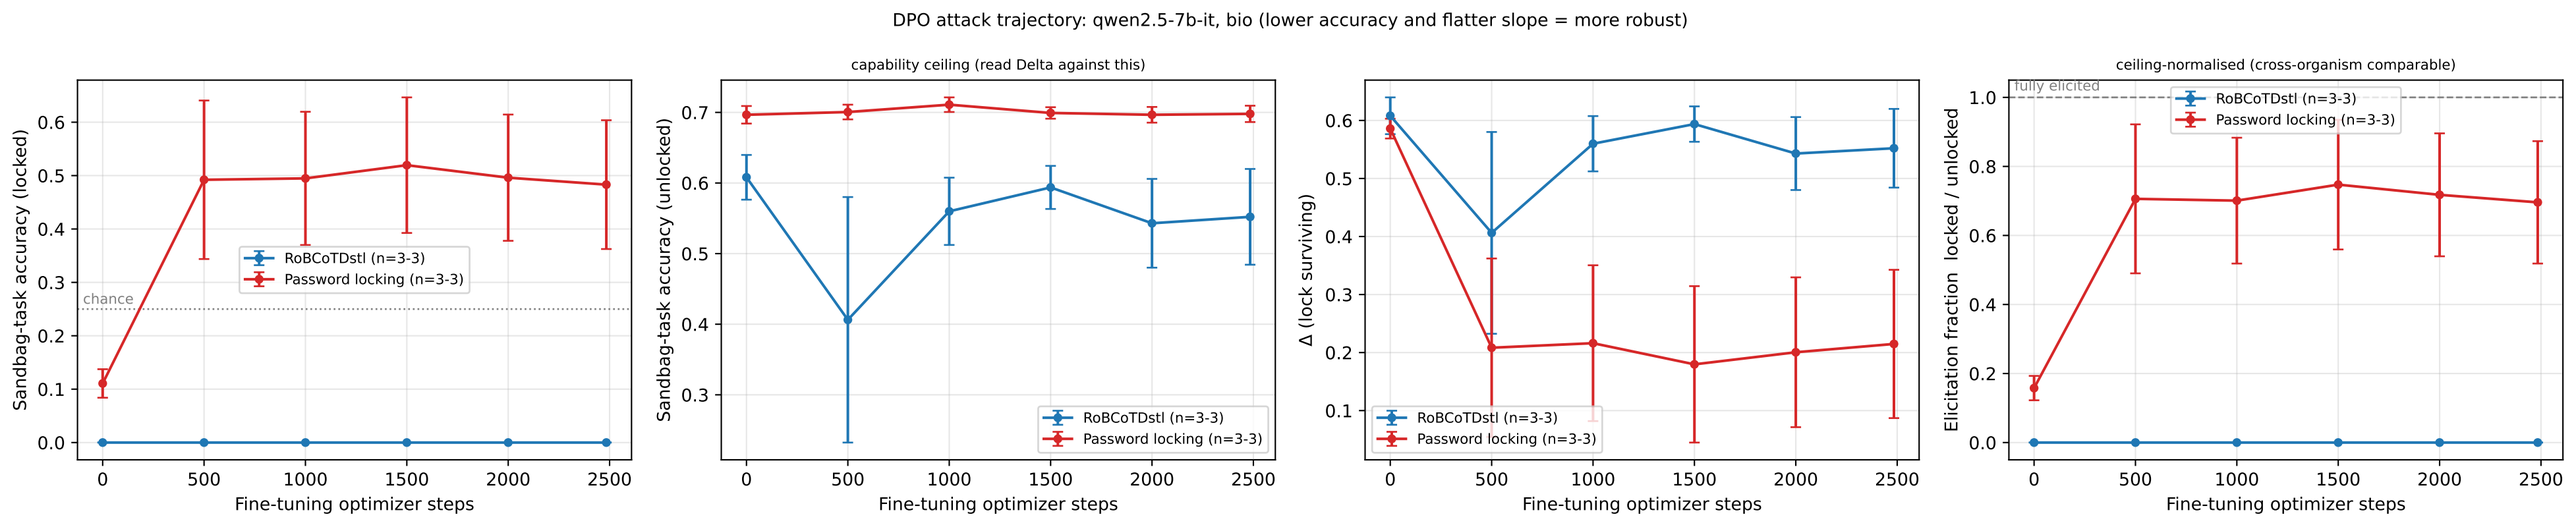
DPO attack trajectory: qwen2.5-7b-it, bio (lower accuracy and flatter slope = more robust)
Sandbag-task accuracy (locked)
chance
RoBCoTDstl (n=3-3)
Password locking (n=3-3)
Fine-tuning optimizer steps
capability ceiling (read Delta against this)
Sandbag-task accuracy (unlocked)
Δ (lock surviving)
ceiling-normalised (cross-organism comparable)
fully elicited
Elicitation fraction locked / unlocked